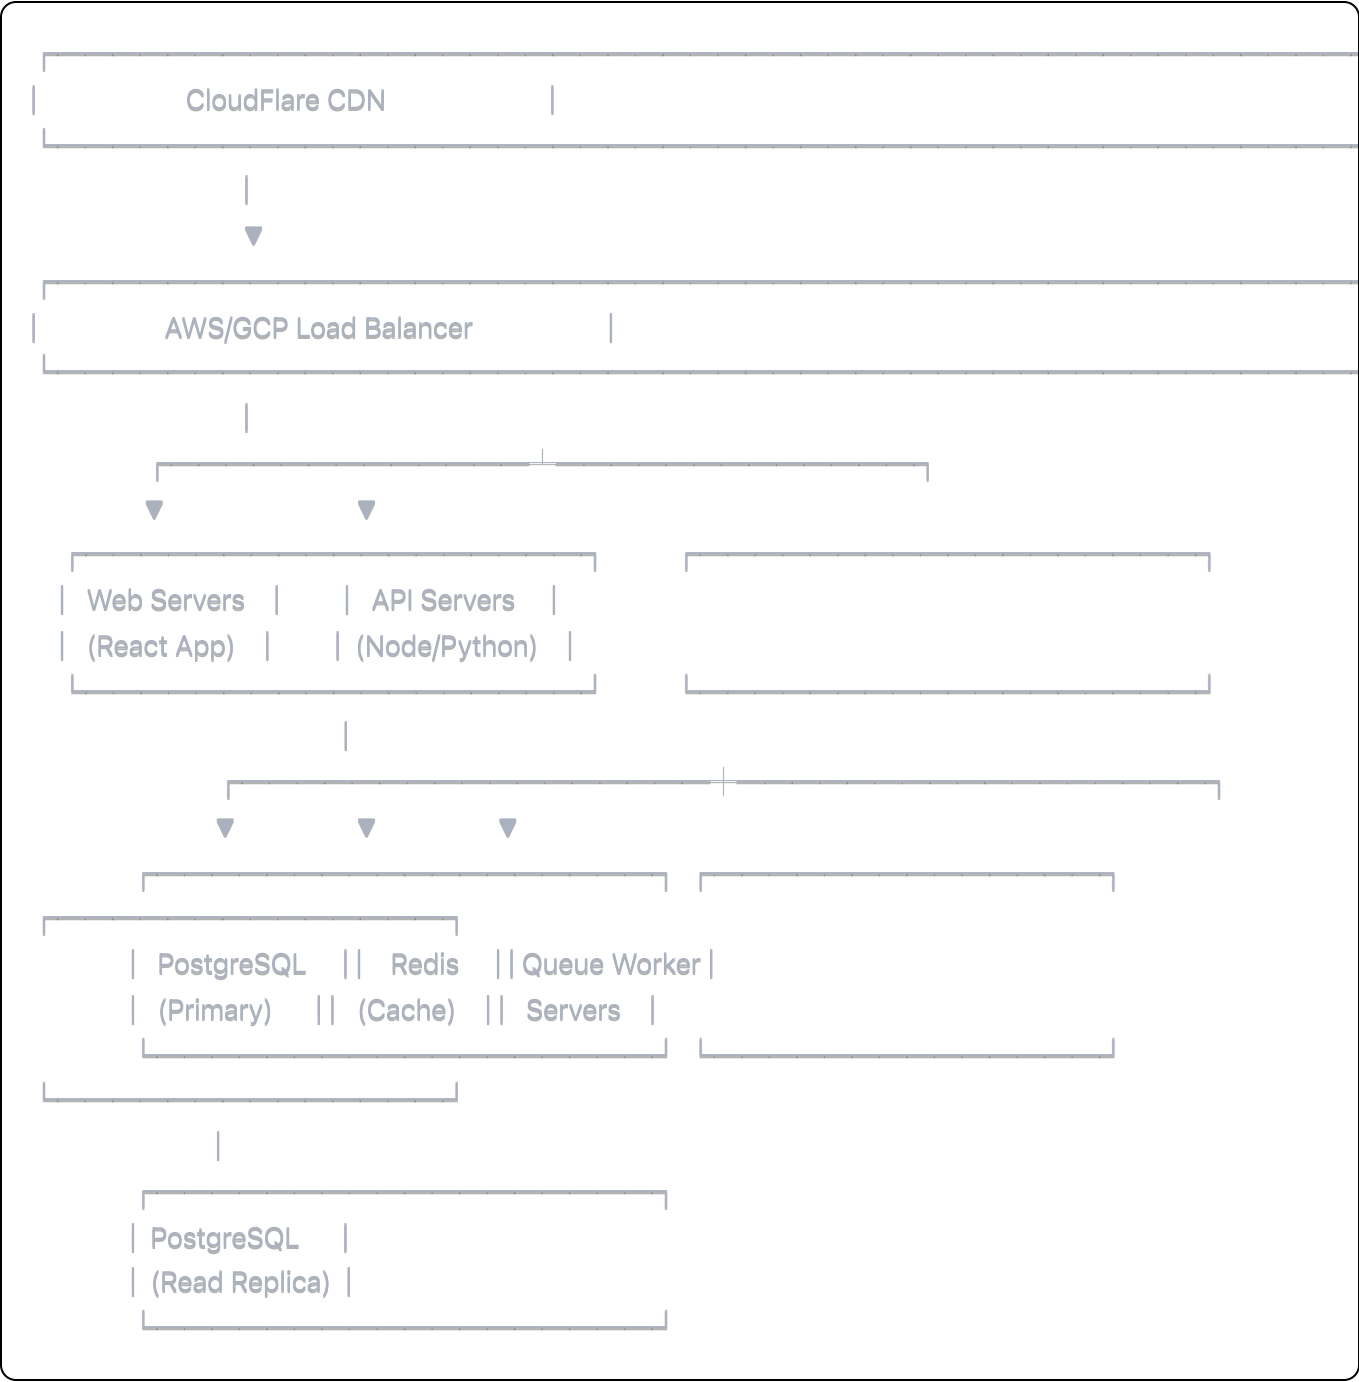


Deployment & Operations Guide - Spotify Follow-Swarm

Production Architecture

Infrastructure Overview



Docker Configuration

Docker Compose (Development)

yaml

```
# docker-compose.yml
```

```
version: '3.8'
```

```
services:
```

```
  postgres:
```

```
    image: postgres:14
```

```
    environment:
```

```
      POSTGRES_DB: spotify_swarm
```

```
      POSTGRES_USER: admin
```

```
      POSTGRES_PASSWORD: ${DB_PASSWORD}
```

```
    volumes:
```

```
      - postgres_data:/var/lib/postgresql/data
```

```
    ports:
```

```
      - "5432:5432"
```

```
  redis:
```

```
    image: redis:7-alpine
```

```
    ports:
```

```
      - "6379:6379"
```

```
    volumes:
```

```
      - redis_data:/data
```

```
  api:
```

```
    build:
```

```
      context: ./backend
```

```
      dockerfile: Dockerfile
```

```
    environment:
```

```
      DATABASE_URL: postgresql://admin:${DB_PASSWORD}@postgres:5432/spotify_swarm
```

```
      REDIS_URL: redis://redis:6379
```

```
      SPOTIFY_CLIENT_ID: ${SPOTIFY_CLIENT_ID}
```

```
      SPOTIFY_CLIENT_SECRET: ${SPOTIFY_CLIENT_SECRET}
```

```
    depends_on:
```

```
      - postgres
```

```
      - redis
```

```
    ports:
```

```
      - "3001:3001"
```

```
    volumes:
```

```
      - ./backend:/app
```

```
      - /app/node_modules
```

```
  worker:
```

```
build:
  context: ./backend
  dockerfile: Dockerfile.worker
environment:
  DATABASE_URL: postgresql://admin:${DB_PASSWORD}@postgres:5432/spotify_swarm
  REDIS_URL: redis://redis:6379
depends_on:
  - postgres
  - redis
  - api
volumes:
  - ./backend:/app
  - /app/node_modules

frontend:
  build:
    context: ./frontend
    dockerfile: Dockerfile
  environment:
    REACT_APP_API_URL: http://localhost:3001
  ports:
    - "3000:3000"
  volumes:
    - ./frontend:/app
    - /app/node_modules

volumes:
  postgres_data:
  redis_data:
```

Production Dockerfiles

Backend API Dockerfile

dockerfile

```
# backend/Dockerfile
FROM node:18-alpine AS builder

WORKDIR /app
COPY package*.json ./
RUN npm ci --only=production

FROM node:18-alpine

WORKDIR /app
COPY --from=builder /app/node_modules ./node_modules
COPY . .

# Run as non-root user
RUN addgroup -g 1001 -S nodejs && \
    adduser -S nodejs -u 1001
USER nodejs

EXPOSE 3001
CMD ["node", "server.js"]
```

Worker Dockerfile

```
dockerfile

# backend/Dockerfile.worker
FROM node:18-alpine

WORKDIR /app
COPY package*.json ./
RUN npm ci --only=production
COPY . .

# Run as non-root user
RUN addgroup -g 1001 -S nodejs && \
    adduser -S nodejs -u 1001
USER nodejs

CMD ["node", "worker.js"]
```

Kubernetes Deployment

API Deployment

yaml

```
# k8s/api-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: api
  template:
    metadata:
      labels:
        app: api
    spec:
      containers:
        - name: api
          image: spotify-swarm-api:latest
          ports:
            - containerPort: 3001
          env:
            - name: DATABASE_URL
              valueFrom:
                secretKeyRef:
                  name: app-secrets
                  key: database-url
            - name: REDIS_URL
              valueFrom:
                secretKeyRef:
                  name: app-secrets
                  key: redis-url
      resources:
        requests:
          memory: "256Mi"
          cpu: "250m"
        limits:
          memory: "512Mi"
          cpu: "500m"
      livenessProbe:
        httpGet:
          path: /health
```

```
  port: 3001
  initialDelaySeconds: 30
  periodSeconds: 10
readinessProbe:
  httpGet:
    path: /ready
    port: 3001
  initialDelaySeconds: 5
  periodSeconds: 5
```

Worker Deployment

yaml


```
# k8s/worker-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: worker
  template:
    metadata:
      labels:
        app: worker
    spec:
      containers:
        - name: worker
          image: spotify-swarm-worker:latest
          env:
            - name: DATABASE_URL
              valueFrom:
                secretKeyRef:
                  name: app-secrets
                  key: database-url
            - name: REDIS_URL
              valueFrom:
                secretKeyRef:
                  name: app-secrets
                  key: redis-url
          resources:
            requests:
              memory: "512Mi"
              cpu: "500m"
            limits:
              memory: "1Gi"
              cpu: "1000m"
```

Environment Configuration

Production Environment Variables

bash

.env.production

Application

NODE_ENV=production

APP_URL=https://spotifyswarm.com

API_URL=https://api.spotifyswarm.com

Database

DATABASE_URL=postgresql://user:pass@db.amazonaws.com:5432/spotify_swarm

DATABASE_POOL_MIN=2

DATABASE_POOL_MAX=10

Redis

REDIS_URL=redis://redis.amazonaws.com:6379

REDIS_PASSWORD=your_redis_password

Spotify API

SPOTIFY_CLIENT_ID=your_client_id

SPOTIFY_CLIENT_SECRET=your_client_secret

SPOTIFY_REDIRECT_URI=https://spotifyswarm.com/auth/callback

Security

JWT_SECRET=your_jwt_secret_key

ENCRYPTION_KEY=your_encryption_key

SESSION_SECRET=your_session_secret

Monitoring

SENTRY_DSN=https://key@sentry.io/project

NEW_RELIC_LICENSE_KEY=your_license_key

DATADOG_API_KEY=your_api_key

Email

SENDGRID_API_KEY=your_sendgrid_key

EMAIL_FROM=noreply@spotifyswarm.com

Payments

STRIPE_SECRET_KEY=sk_live_your_key

STRIPE_WEBHOOK_SECRET=whsec_your_secret

CI/CD Pipeline

GitHub Actions Workflow

yaml

```
# .github/workflows/deploy.yml
```

```
name: Deploy to Production
```

```
on:
```

```
  push:
```

```
    branches: [main]
```

```
  pull_request:
```

```
    branches: [main]
```

```
jobs:
```

```
  test:
```

```
    runs-on: ubuntu-latest
```

```
    steps:
```

```
      - uses: actions/checkout@v2
```

```
      - name: Setup Node.js
```

```
        uses: actions/setup-node@v2
```

```
        with:
```

```
          node-version: '18'
```

```
      - name: Install dependencies
```

```
        run: npm ci
```

```
      - name: Run tests
```

```
        run: npm test
```

```
      - name: Run linting
```

```
        run: npm run lint
```

```
  build:
```

```
    needs: test
```

```
    runs-on: ubuntu-latest
```

```
    steps:
```

```
      - uses: actions/checkout@v2
```

```
      - name: Build Docker images
```

```
        run: |
```

```
          docker build -t spotify-swarm-api:${{ github.sha }} ./backend
```

```
          docker build -t spotify-swarm-worker:${{ github.sha }} ./backend -f Dockerfile.worker
```

```
          docker build -t spotify-swarm-frontend:${{ github.sha }} ./frontend
```

```
- name: Push to Registry
```

```
env:
```

```
  DOCKER_REGISTRY: ${ secrets.DOCKER_REGISTRY }
```

```
run: |
```

```
  echo ${ secrets.DOCKER_PASSWORD } | docker login -u ${ secrets.DOCKER_USERNAME } --pass
```

```
  docker push spotify-swarm-api:${ github.sha }
```

```
  docker push spotify-swarm-worker:${ github.sha }
```

```
  docker push spotify-swarm-frontend:${ github.sha }
```

```
deploy:
```

```
  needs: build
```

```
  runs-on: ubuntu-latest
```

```
  if: github.ref == 'refs/heads/main'
```

```
  steps:
```

```
- name: Deploy to Kubernetes
```

```
env:
```

```
  KUBE_CONFIG: ${ secrets.KUBE_CONFIG }
```

```
run: |
```

```
  echo "$KUBE_CONFIG" | base64 -d > kubeconfig
```

```
  export KUBECONFIG=kubeconfig
```

```
  kubectl set image deployment/api-deployment api=spotify-swarm-api:${ github.sha }
```

```
  kubectl set image deployment/worker-deployment worker=spotify-swarm-worker:${ github.sha }
```

```
  kubectl rollout status deployment/api-deployment
```

```
  kubectl rollout status deployment/worker-deployment
```

Monitoring & Alerting

Health Check Endpoints

```
javascript
```

```

// health-checks.js
app.get('/health', (req, res) => {
  res.status(200).json({ status: 'healthy' });
});

app.get('/ready', async (req, res) => {
  try {
    // Check database
    await db.query('SELECT 1');

    // Check Redis
    await redis.ping();

    // Check Spotify API
    const token = await getSystemToken();
    if (!token) throw new Error('No system token');

    res.status(200).json({
      status: 'ready',
      services: {
        database: 'connected',
        redis: 'connected',
        spotify: 'authenticated'
      }
    });
  } catch (error) {
    res.status(503).json({
      status: 'not ready',
      error: error.message
    });
  }
});

app.get('/metrics', async (req, res) => {
  // Prometheus metrics endpoint
  res.set('Content-Type', register.contentType);
  res.end(await register.metrics());
});

```

Monitoring Stack Configuration

yaml

```
# monitoring/prometheus.yml
global:
  scrape_interval: 15s

scrape_configs:
  - job_name: 'api'
    kubernetes_sd_configs:
      - role: pod
    relabel_configs:
      - source_labels: [__meta_kubernetes_pod_label_app]
        action: keep
        regex: api
      - source_labels: [__meta_kubernetes_pod_ip]
        target_label: __address__
        replacement: $1:3001

  - job_name: 'worker'
    kubernetes_sd_configs:
      - role: pod
    relabel_configs:
      - source_labels: [__meta_kubernetes_pod_label_app]
        action: keep
        regex: worker
```

Alert Rules

yaml

```
# monitoring/alerts.yml
```

```
groups:
```

```
- name: spotify_swarm
```

```
rules:
```

```
- alert: HighErrorRate
```

```
  expr: rate(http_requests_total{status=~"5.."}[5m]) > 0.05
```

```
  for: 5m
```

```
  annotations:
```

```
    summary: High error rate detected
```

```
    description: "Error rate is {{ $value }} errors per second"
```

```
- alert: SlowAPIResponse
```

```
  expr: histogram_quantile(0.95, rate(http_request_duration_seconds_bucket[5m])) > 1
```

```
  for: 10m
```

```
  annotations:
```

```
    summary: API response time is slow
```

```
    description: "95th percentile response time is {{ $value }} seconds"
```

```
- alert: QueueBacklog
```

```
  expr: queue_size{status="pending"} > 1000
```

```
  for: 15m
```

```
  annotations:
```

```
    summary: Large queue backlog
```

```
    description: "{{ $value }} jobs pending in queue"
```

```
- alert: RateLimitNearLimit
```

```
  expr: rate(spotify_api_calls_total[1h]) > 2700
```

```
  for: 5m
```

```
  annotations:
```

```
    summary: Approaching Spotify rate limit
```

```
    description: "Current rate: {{ $value }} calls per hour"
```

Database Management

Backup Strategy

```
bash
```



```
#!/bin/bash
# backup.sh

# Daily backup script
BACKUP_DIR="/backups"
DATE=$(date +%Y%m%d_%H%M%S)
DB_NAME="spotify_swarm"

# Create backup
pg_dump $DATABASE_URL > $BACKUP_DIR/backup_$DATE.sql

# Compress backup
gzip $BACKUP_DIR/backup_$DATE.sql

# Upload to S3
aws s3 cp $BACKUP_DIR/backup_$DATE.sql.gz s3://spotify-swarm-backups/

# Delete local backups older than 7 days
find $BACKUP_DIR -name "backup_*.sql.gz" -mtime +7 -delete

# Delete S3 backups older than 30 days
aws s3 ls s3://spotify-swarm-backups/ | \
while read -r line; do
    createDate=$(echo $line | awk {'print $1' '$2'})
    createDate=$(date -d"$createDate" +%s)
    olderThan=$(date -d"30 days ago" +%s)
    if [[ $createDate -lt $olderThan ]]; then
        fileName=$(echo $line | awk {'print $4'})
        aws s3 rm s3://spotify-swarm-backups/$fileName
    fi
done
```

Migration Strategy

javascript

```
// migrations/run.js
const { Pool } = require('pg');
const fs = require('fs');
const path = require('path');

async function runMigrations() {
  const pool = new Pool({ connectionString: process.env.DATABASE_URL });

  // Create migrations table
  await pool.query(`
    CREATE TABLE IF NOT EXISTS migrations (
      id SERIAL PRIMARY KEY,
      filename VARCHAR(255) UNIQUE NOT NULL,
      executed_at TIMESTAMP DEFAULT NOW()
    )
  `);

  // Get pending migrations
  const files = fs.readdirSync(path.join(__dirname, 'sql'));
  const executed = await pool.query('SELECT filename FROM migrations');
  const executedFiles = executed.rows.map(r => r.filename);

  const pending = files.filter(f => !executedFiles.includes(f)).sort();

  // Run pending migrations
  for (const file of pending) {
    console.log(`Running migration: ${file}`);
    const sql = fs.readFileSync(path.join(__dirname, 'sql', file), 'utf8');

    await pool.query('BEGIN');
    try {
      await pool.query(sql);
      await pool.query('INSERT INTO migrations (filename) VALUES ($1)', [file]);
      await pool.query('COMMIT');
      console.log(`✓ Migration ${file} completed`);
    } catch (error) {
      await pool.query('ROLLBACK');
      console.error(`✗ Migration ${file} failed:`, error);
      process.exit(1);
    }
  }
}
```

```
await pool.end();  
}  
  
runMigrations();
```

Security Checklist

Pre-Deployment Security Audit

- ☐ All dependencies updated to latest secure versions
- ☐ Environment variables properly secured
- ☐ Database connections use SSL
- ☐ API endpoints have rate limiting
- ☐ Input validation on all user inputs
- ☐ XSS protection headers configured
- ☐ CSRF tokens implemented
- ☐ Content Security Policy configured
- ☐ HTTPS enforced everywhere
- ☐ Secrets rotated from development

Runtime Security Monitoring

```
javascript
```

```
// security-monitor.js
const helmet = require('helmet');
const rateLimit = require('express-rate-limit');
const mongoSanitize = require('express-mongo-sanitize');

// Security middleware
app.use(helmet({
  contentSecurityPolicy: {
    directives: {
      defaultSrc: ["'self'"],
      scriptSrc: ["'self'", "'unsafe-inline'"],
      styleSrc: ["'self'", "'unsafe-inline'"],
      imgSrc: ["'self'", "data:", "https:"],
    },
  },
}));

// Rate limiting
const limiter = rateLimit({
  windowMs: 15 * 60 * 1000, // 15 minutes
  max: 100 // limit each IP to 100 requests per windowMs
});
app.use('/api/', limiter);

// Input sanitization
app.use(mongoSanitize());

// Security logging
app.use((req, res, next) => {
  if (req.path.includes('admin') || req.method !== 'GET') {
    logger.info({
      type: 'security',
      method: req.method,
      path: req.path,
      ip: req.ip,
      user: req.user?.id
    });
  }
  next();
});
```

Performance Optimization

Caching Strategy

javascript

```

// cache-manager.js
class CacheManager {
  constructor(redis) {
    this.redis = redis;
    this.ttl = {
      user: 3600,    // 1 hour
      follows: 300,  // 5 minutes
      analytics: 1800, // 30 minutes
      static: 86400  // 24 hours
    };
  }

  async get(key, fetcher, ttl) {
    // Try cache first
    const cached = await this.redis.get(key);
    if (cached) {
      return JSON.parse(cached);
    }

    // Fetch and cache
    const data = await fetcher();
    await this.redis.setex(
      key,
      ttl || this.ttl.static,
      JSON.stringify(data)
    );

    return data;
  }

  async invalidate(pattern) {
    const keys = await this.redis.keys(pattern);
    if (keys.length > 0) {
      await this.redis.del(...keys);
    }
  }
}

```

Database Optimization

sql

```
-- Indexes for performance
CREATE INDEX idx_follows_follower_status ON follows(follower_id, status);
CREATE INDEX idx_follows_followed_status ON follows(followed_id, status);
CREATE INDEX idx_queue_jobs_status ON queue_jobs(status, scheduled_for);
CREATE INDEX idx_analytics_user_created ON analytics(user_id, created_at);

-- Partitioning for analytics table
CREATE TABLE analytics_2024_01 PARTITION OF analytics
FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');

-- Query optimization views
CREATE MATERIALIZED VIEW user_stats AS
SELECT
  u.id,
  COUNT(DISTINCT f1.followed_id) as following_count,
  COUNT(DISTINCT f2.follower_id) as follower_count,
  MAX(f1.completed_at) as last_follow_at
FROM users u
LEFT JOIN follows f1 ON f1.follower_id = u.id AND f1.status = 'completed'
LEFT JOIN follows f2 ON f2.followed_id = u.id AND f2.status = 'completed'
GROUP BY u.id;

-- Refresh materialized view periodically
CREATE OR REPLACE FUNCTION refresh_user_stats()
RETURNS void AS $$
BEGIN
  REFRESH MATERIALIZED VIEW CONCURRENTLY user_stats;
END;
$$ LANGUAGE plpgsql;
```

Disaster Recovery

Rollback Procedure

```
bash
```

```
#!/bin/bash
# rollback.sh

PREVIOUS_VERSION=$1

if [ -z "$PREVIOUS_VERSION" ]; then
    echo "Usage: ./rollback.sh <version>"
    exit 1
fi

echo "Rolling back to version $PREVIOUS_VERSION"

# Rollback Kubernetes deployments
kubectl set image deployment/api-deployment api=spotify-swarm-api:$PREVIOUS_VERSION
kubectl set image deployment/worker-deployment worker=spotify-swarm-worker:$PREVIOUS_VERSION

# Wait for rollout
kubectl rollout status deployment/api-deployment
kubectl rollout status deployment/worker-deployment

# Clear cache
redis-cli FLUSHALL

echo "Rollback complete"
```

Incident Response Plan

1. **Detection** - Alert triggered via monitoring
2. **Assessment** - Check dashboards and logs
3. **Communication** - Notify team and update status page
4. **Mitigation** - Apply immediate fix or rollback
5. **Resolution** - Deploy permanent fix
6. **Post-mortem** - Document and learn from incident

Scaling Guidelines

Horizontal Scaling Triggers

- CPU usage > 70% for 5 minutes

- Memory usage > 80% for 5 minutes
- Request latency p95 > 1 second
- Queue depth > 5000 jobs

Vertical Scaling Recommendations

- **API Servers:** Start with 2 CPU, 4GB RAM
- **Workers:** Start with 4 CPU, 8GB RAM
- **Database:** Start with db.t3.medium (2 vCPU, 4GB)
- **Redis:** Start with cache.t3.micro (2 vCPU, 0.5GB)